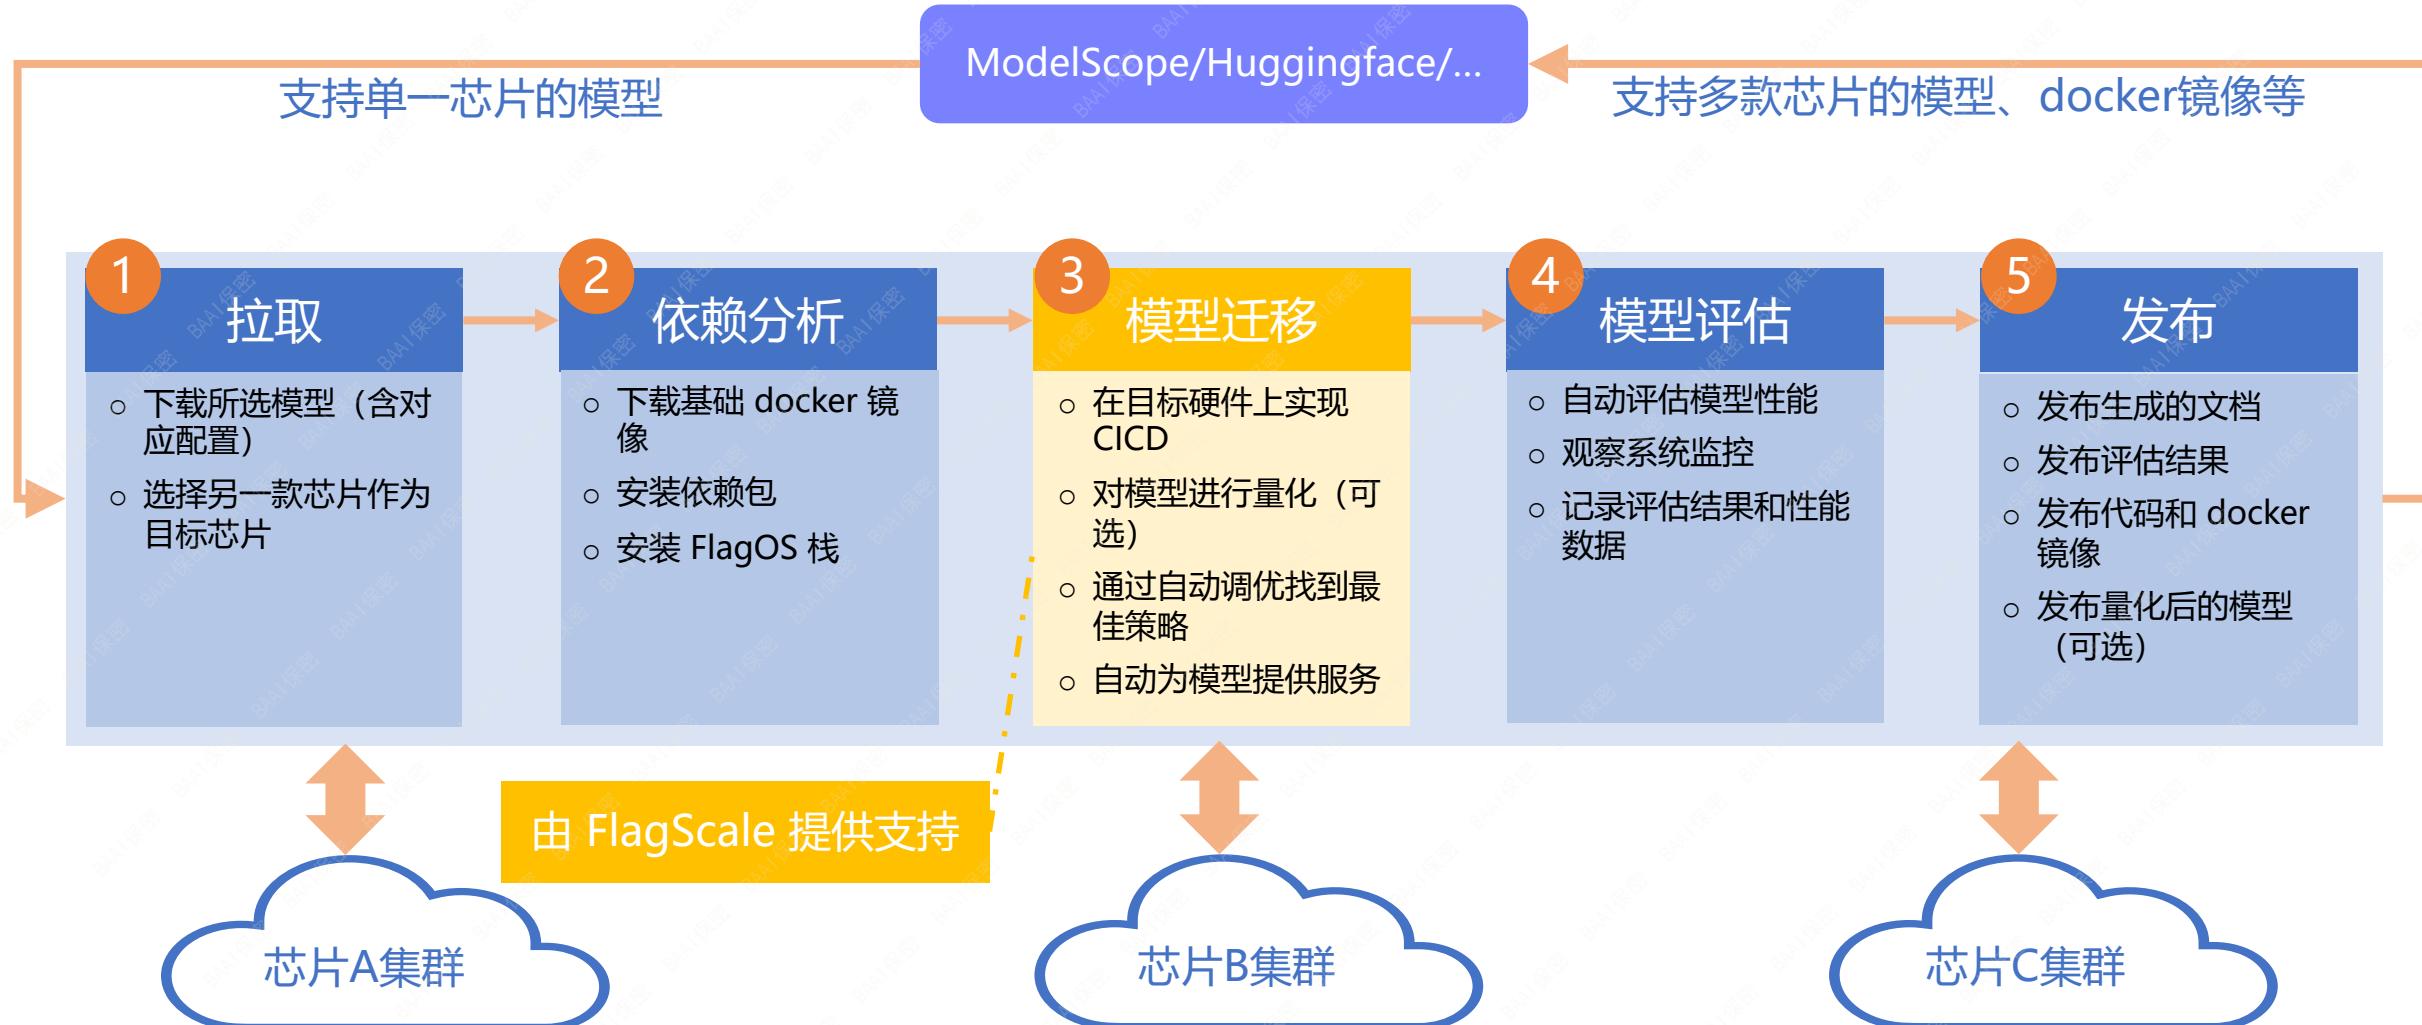
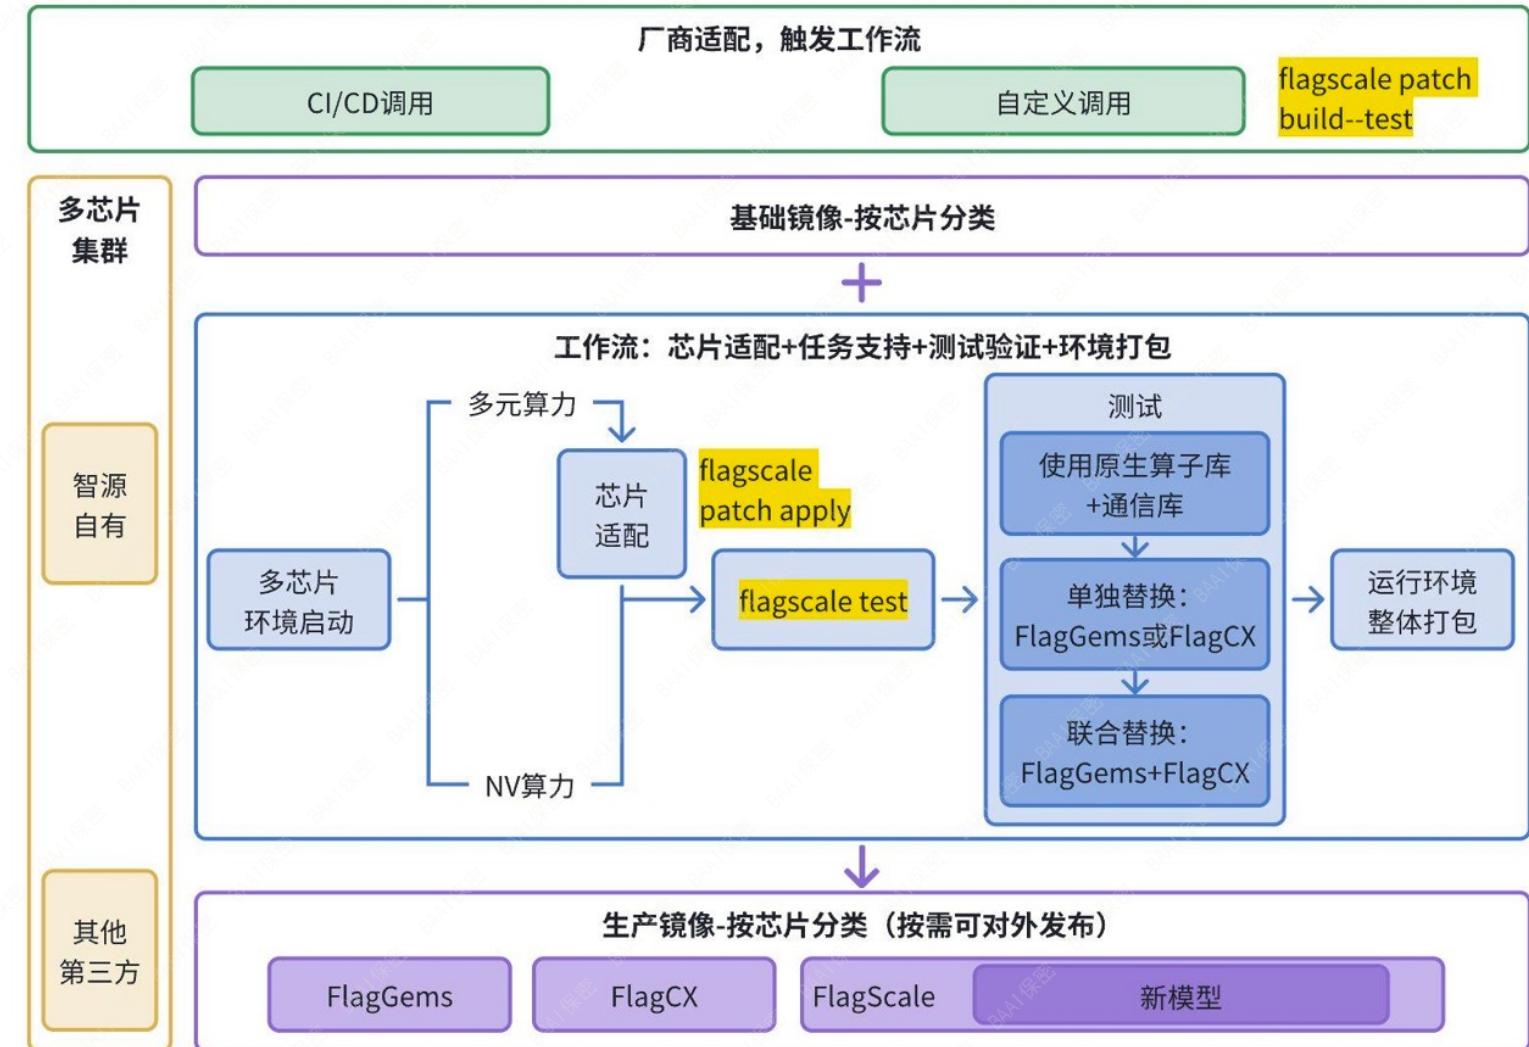


基于 FlagRelease 实现模型到不同芯片上的推理部署



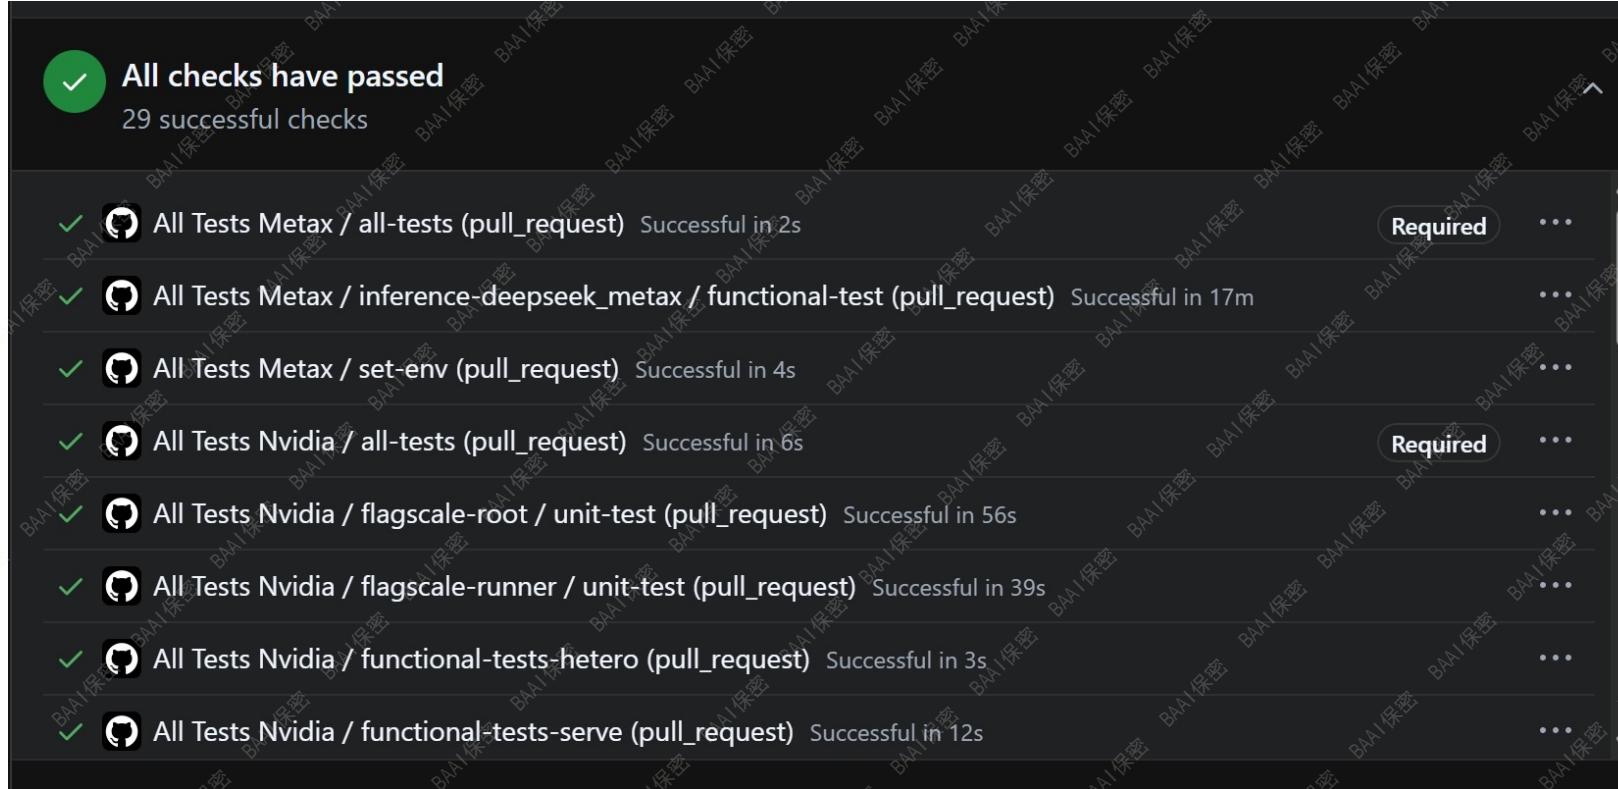
自动化CI/CD持续升级完善，更好支持FlagRelease

- 提交PR时自动触发
- 本地可手动出发
- 集成高性能算子库FlagGems测试
- 未来将继承高性能统一通信库FlagCX测试
- 为FlagRelease新模型的发布提供支撑



自动化CI/CD持续升级完善，更好支持FlagRelease

- 当前已上线沐曦CICD，更多国产芯片CI/CD未来持续上线中



- **下载镜像**

```
docker pull flagorelease-registry.cn-beijing.cr.aliyuncs.com/flagorelease/flagorelease:nv_vllm085_gemscbbf39
```

- **启动容器**

```
docker run --rm --init --detach \ --net=host --uts=host --ipc=host \ --security-opt=seccomp=unconfined \ --privileged=true \ --ulimit stack=67108864 \ --ulimit memlock=-1 \ --ulimit nofile=1048576:1048576 \ --shm-size=32G \ -v /nfs:/nfs \ --gpus all \ --name flagos \ flagorelease-registry.cn-beijing.cr.aliyuncs.com/flagorelease/flagorelease:nv_vllm085_gemscbbf39 \ sleep infinity docker exec -it flagos bash
```

- **安装modelscope并下载模型**

```
pip install modelscope
```

```
modelscope download --model Qwen/Qwen3-4B --local_dir /nfs/Qwen3-4B
```

- **启动推理服务**

```
flagscale serve qwen3
```

<https://www.modelscope.cn/models/FlagRelease/Qwen3-4B-FlagOS-Nvidia>

- **下载镜像**

```
docker pull flagorelease-registry.cn-beijing.cr.aliyuncs.com/flagorelease/flagrelease:flagrelease_metax_qwen3
```

- **启动容器**

```
docker run -it --device=/dev/dri --device=/dev/mxcd --group-add video --name test_flagrelease --  
device=/dev/mem --network=host --security-opt seccomp=unconfined --security-opt apparmor=unconfined --  
shm-size=100gb --ulimit memlock=-1 -v /usr/local/:/usr/local/ -v /nfs:/nfs flagorelease-registry.cn-  
beijing.cr.aliyuncs.com/flagorelease/flagrelease:flagrelease_metax_qwen3 /bin/bash
```

- **安装modelscope并下载模型**

```
pip install modelscope
```

```
modelscope download --model Qwen/Qwen3-4B --local_dir /nfs/Qwen3-4B
```

- **启动推理服务**

```
flagscale serve qwen3
```

<https://www.modelscope.cn/models/FlagRelease/Qwen3-4B-FlagOS-Metax>

谢谢聆听，
欢迎合作与建议。